THE UNIVERSITY OF MELBOURNE

Spatial filtering

Semester 2, 2021

Kris Ehinger

© Copyright The University of Melbourne 2021

https://www12.lunapic.com/editor/?action=sketch

Outline

- Introduction to convolution
- Commonly-used linear filters
- Filters in practice

Learning outcomes

- Explain the convolution and cross-correlation operations
- Identify commonly-used filters and their expected outputs
- Explain practical considerations in implementing filters (efficiency, border handling)

Introduction to convolution

Pixel operator

- Pixel operator: computes an output value at each pixel location, based on the input pixel value

$$g(i,j) = h(f(i,j))$$

Output image g Input image f

- Example: $g(i,j) = 0.5(f(i,j))$

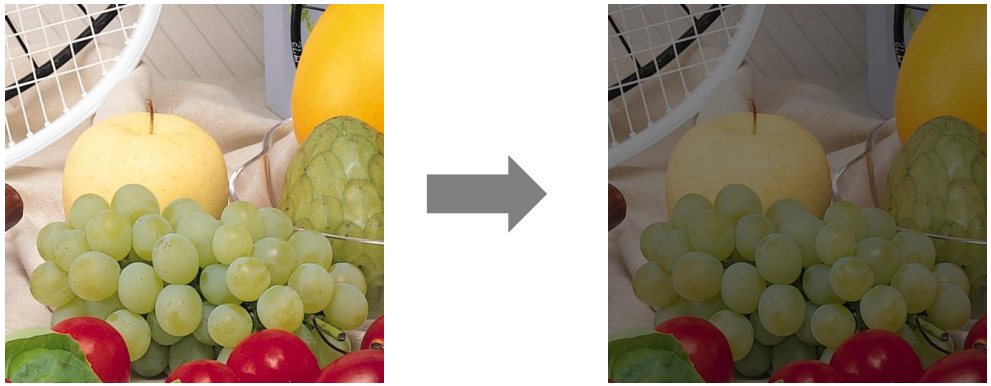

Gamma correction

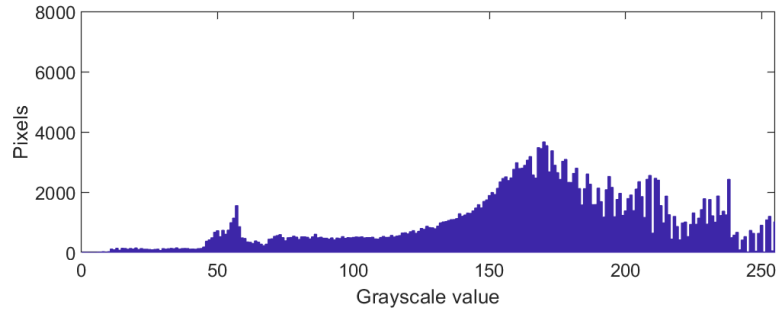

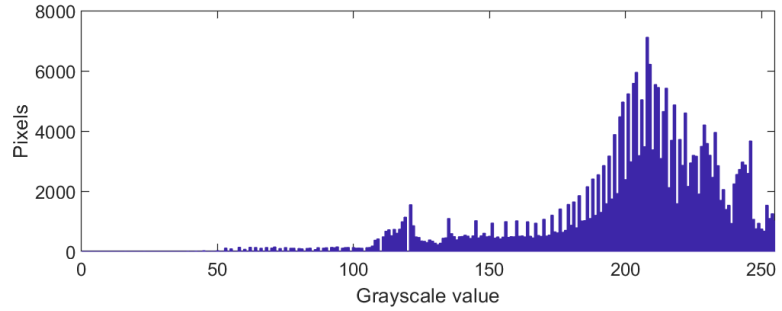

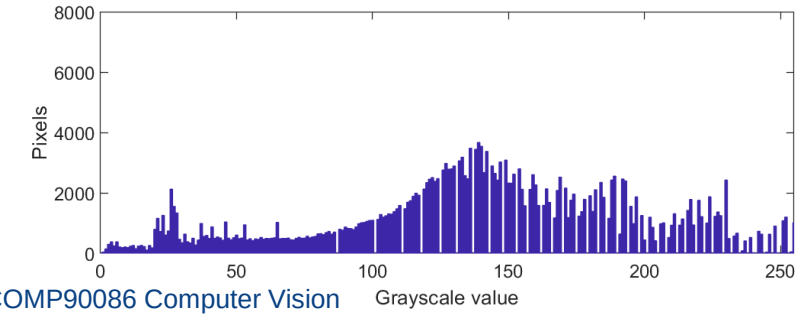

Local operator

- Local operator: computes an output value at each pixel location, based on a neighbourhood of pixels around the input pixel
- Example: sharpening filter

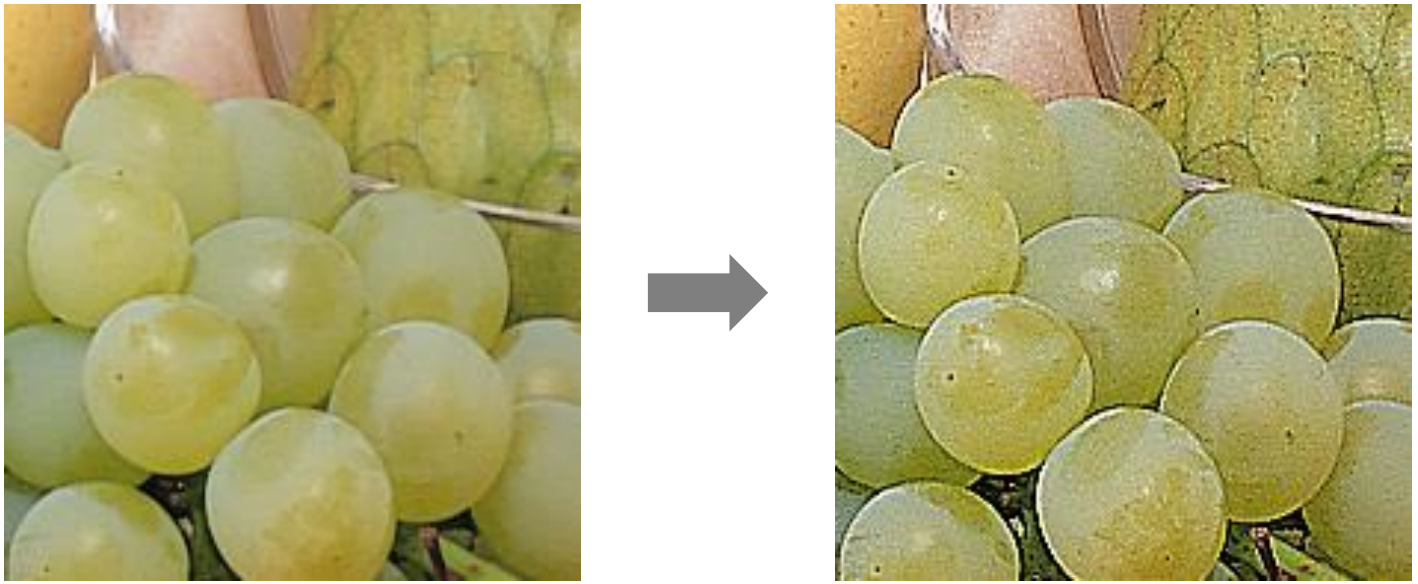

Linear filtering

- Output pixel's value is a weighted sum of a neighbourhood around the input pixel

$$g(i,j) = h(u,v) \otimes f(i,j)$$

Output image g Kernel h Input image f

Cross-correlation convolution

$$g(i,j) = \sum_{u,v} f(i+u, j+v) h(u,v)$$

Linear filtering

- Output pixel's value is a weighted sum of a neighbourhood around the input pixel

$$g(i,j) = h(u,v) * f(i,j)$$

Output image g Kernel h Input image f

Convolution operator

$$g(i,j) = \sum_{u,v} f(i-u, j-v)h(u,v)$$

Convolution

- Consider a 3x4 image and 2x2 kernel

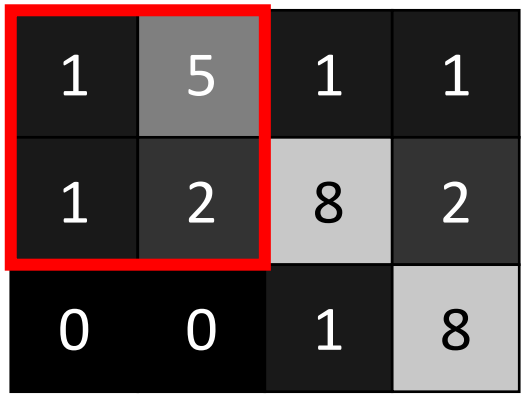

Image

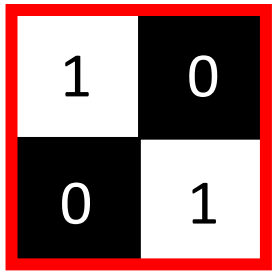

Kernel

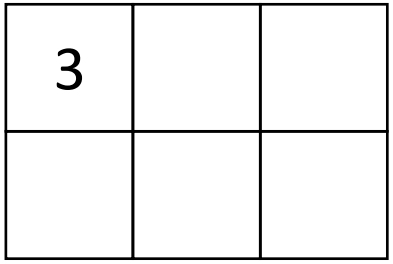

Convolution output

Convolution 1:

1x1	5x0
1x0	2x1

Convolution

- Consider a 3x4 image and 2x2 kernel

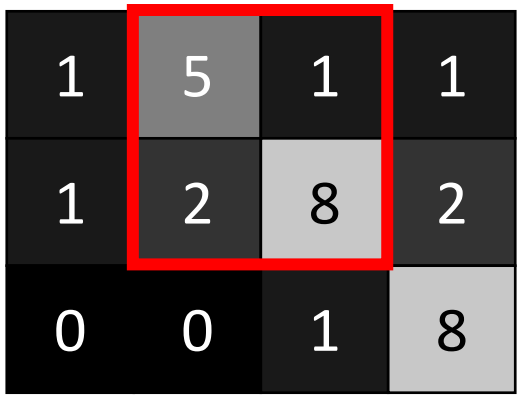

Image

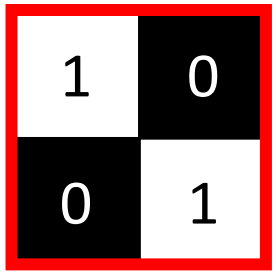

Kernel

3	13	

Convolution output

Convolution 2:

5x1	1x0
2x0	8x1

Convolution

- Consider a 3x4 image and 2x2 kernel

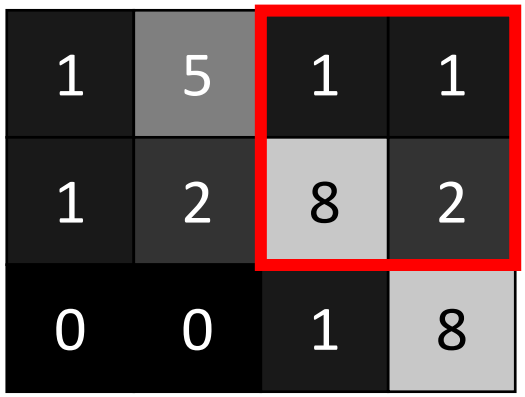

Image

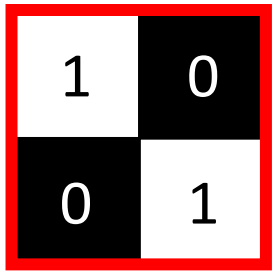

Kernel

3	13	3

Convolution output

Convolution 3:

1x1	1x0
8x0	2x1

Convolution

- Consider a 3x4 image and 2x2 kernel

Image

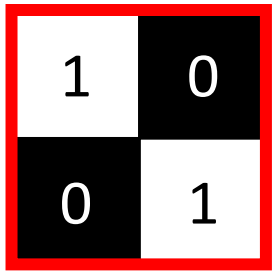

Kernel

3	13	3
1		

Convolution output

Convolution

- Consider a 3x4 image and 2x2 kernel

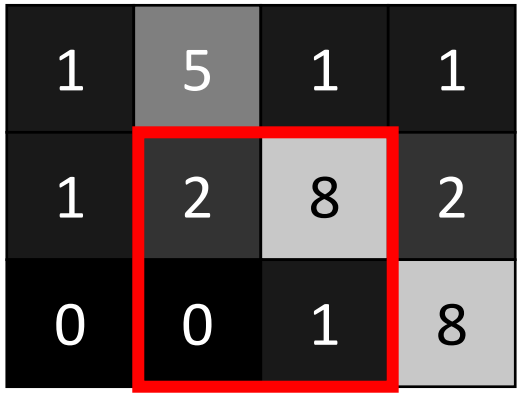

Image

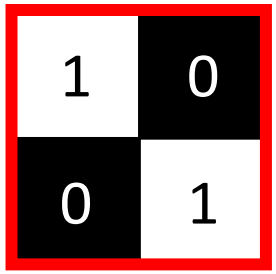

Kernel

3	13	3
1	3	

Convolution output

Convolution

- Consider a 3x4 image and 2x2 kernel

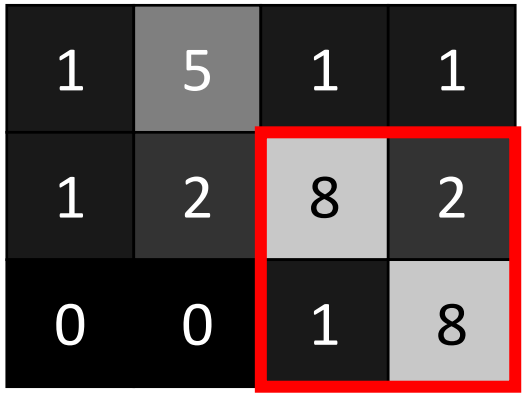

Image

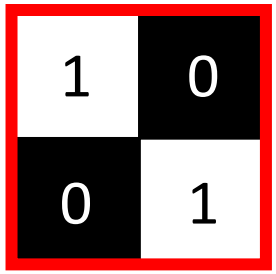

Kernel

3	13	3
1	3	16

Convolution output

Convolution in colour

- Consider a 3x4x3 image and 2x2x3 kernel

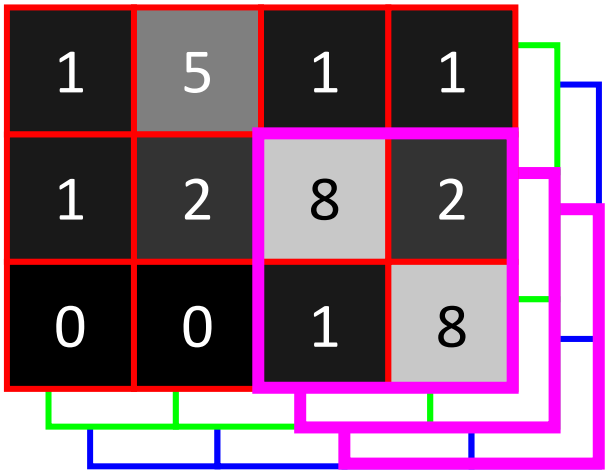

Image

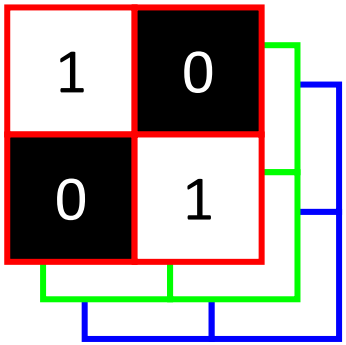

Kernel

Convolution output?

Cross-correlation vs. convolution

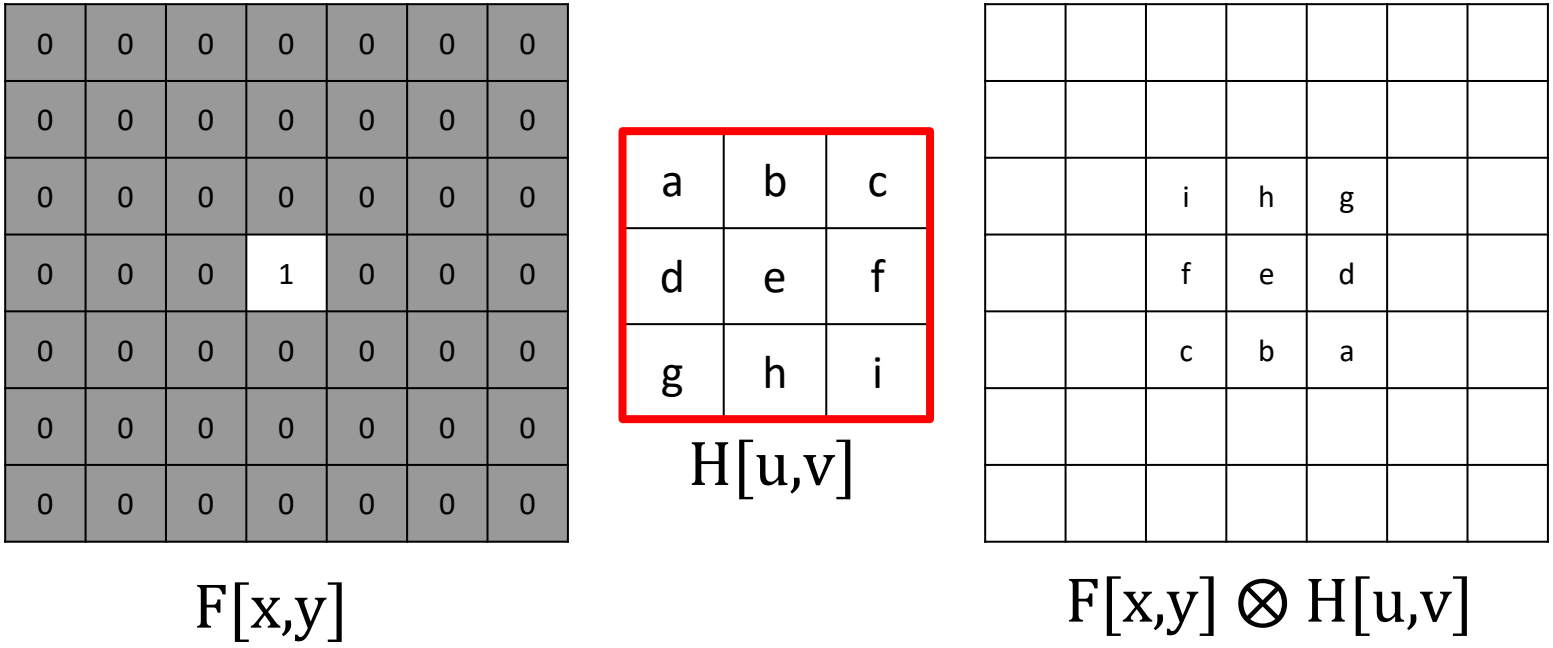

Cross-correlation: overlay filter on image

Image: S. Seitz

Cross-correlation vs. convolution

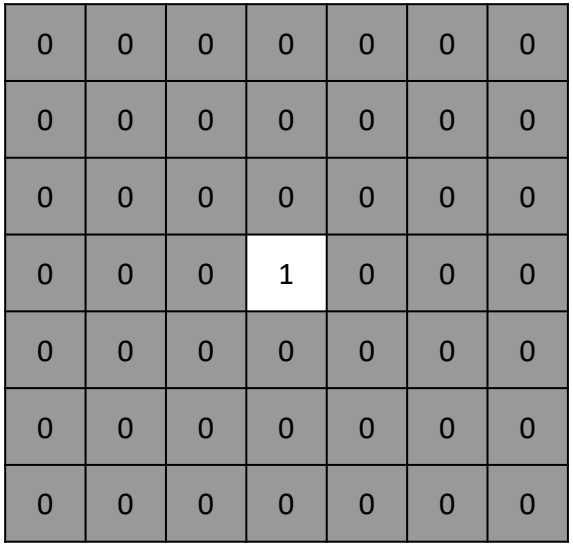

F[x,y]

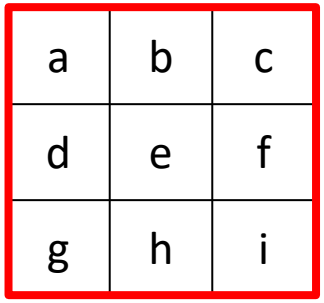

H[u,v]

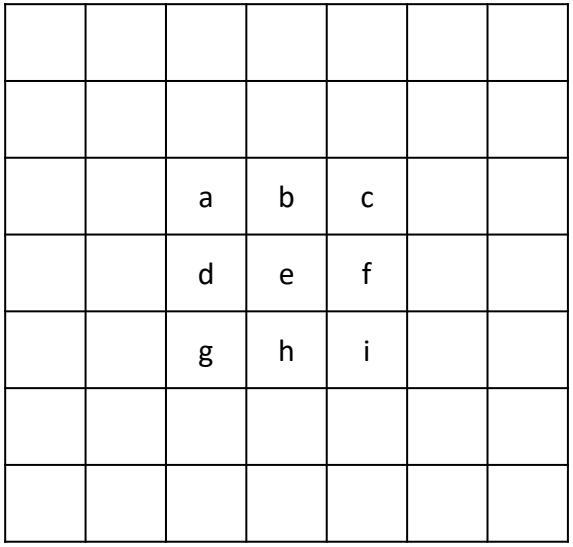

F[x,y] ∗ H[u,v]

Convolution: flip filter horizontally and vertically

Image: S. Seitz

Summary

- Pixel operator: transform pixel based on its value
- Local operator: transform pixel based on its neighbours
- Convolution (and cross-correlation): operations that apply a linear filter to an image

Common filters

Practice with linear filters

Original

0	0	0
0	1	0
0	0	0

?

Slide: D. Lowe

Practice with linear filters

Original

0	0	0
0	1	0
0	0	0

Filtered
(no change)

Slide: D. Lowe

Practice with linear filters

Original

0	0	0
0	0	1
0	0	0

?

Slide: D. Lowe

Practice with linear filters

Original

0	0	0
0	0	1
0	0	0

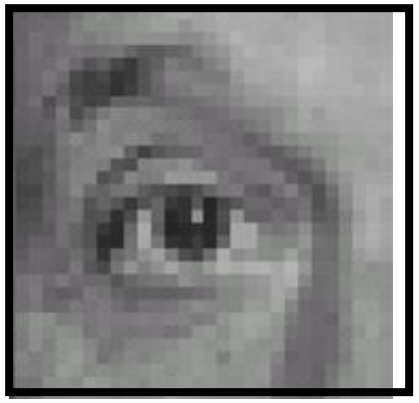

Shifted left
By 1 pixel

Slide: D. Lowe

Practice with linear filters

Original

$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{bmatrix} - \frac{1}{9}\begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} \quad ?$$

(Note that filter sums to 1)

Slide: D. Lowe

Practice with linear filters

Original

$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{bmatrix} - \frac{1}{9}\begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

Sharpening filter

- Accentuates differences with local average

Slide: D. Lowe

Common filters: Gaussian

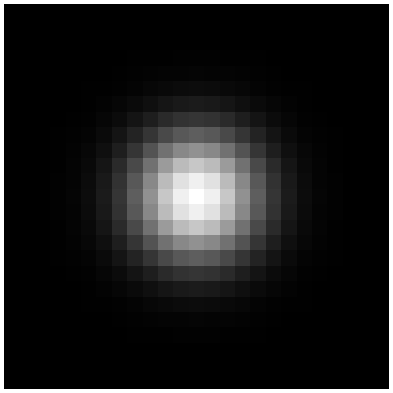

$$G_\sigma = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

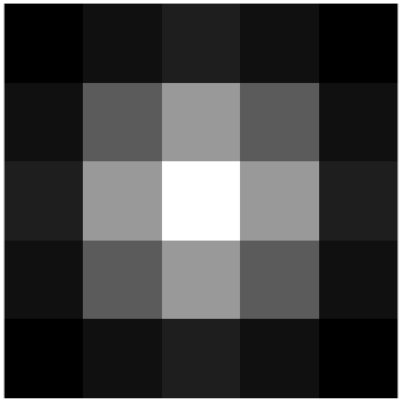

Gaussian kernel
Kernel size: 5 x 5 px
σ = 1

Blur filters

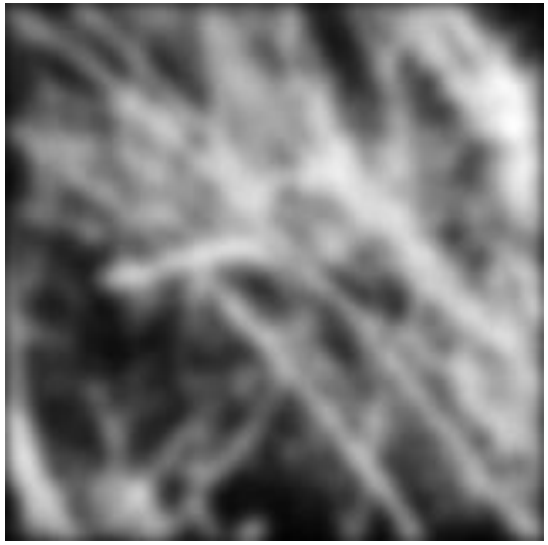

Gaussian kernel

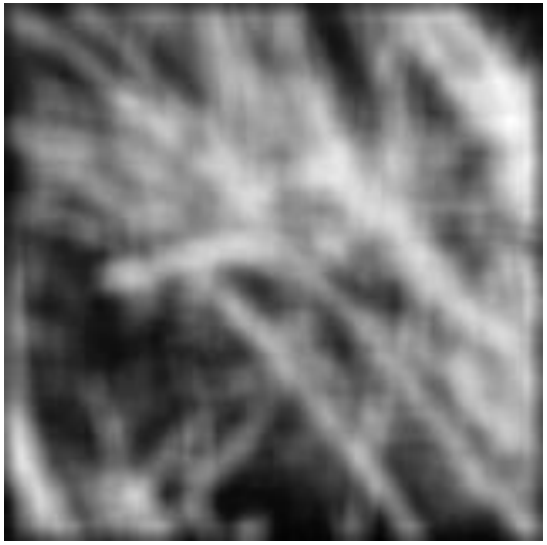

Average kernel

Common filters: Sobel

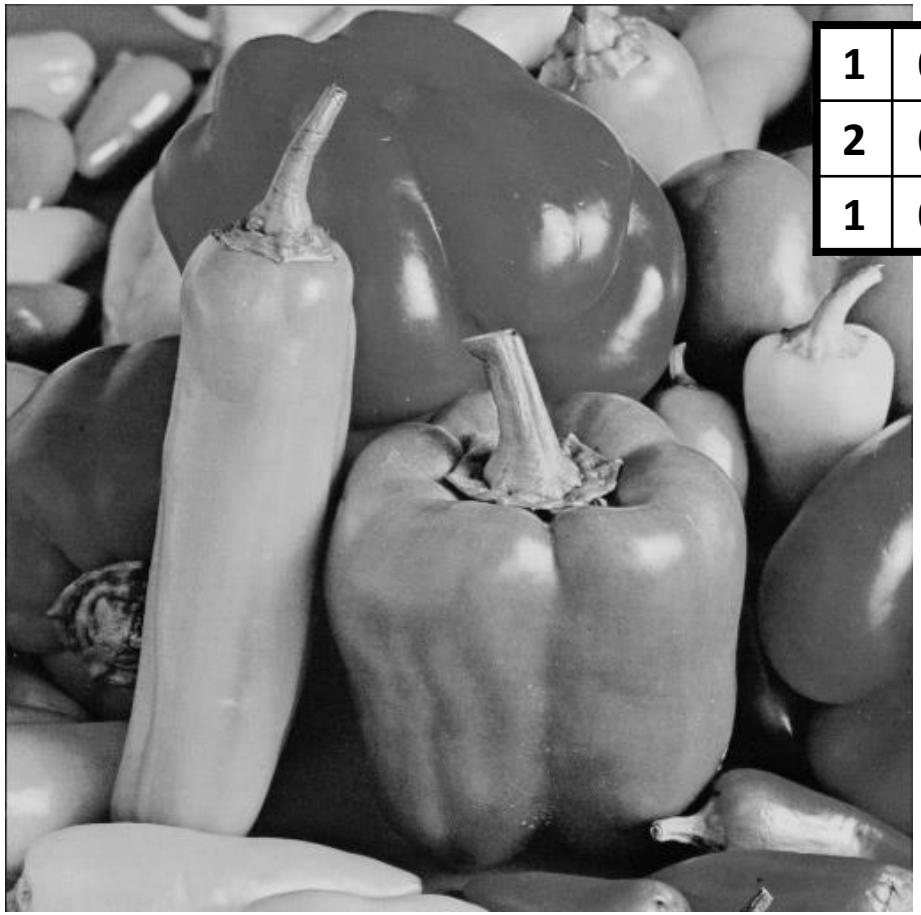

1	0	-1
2	0	-2
1	0	-1

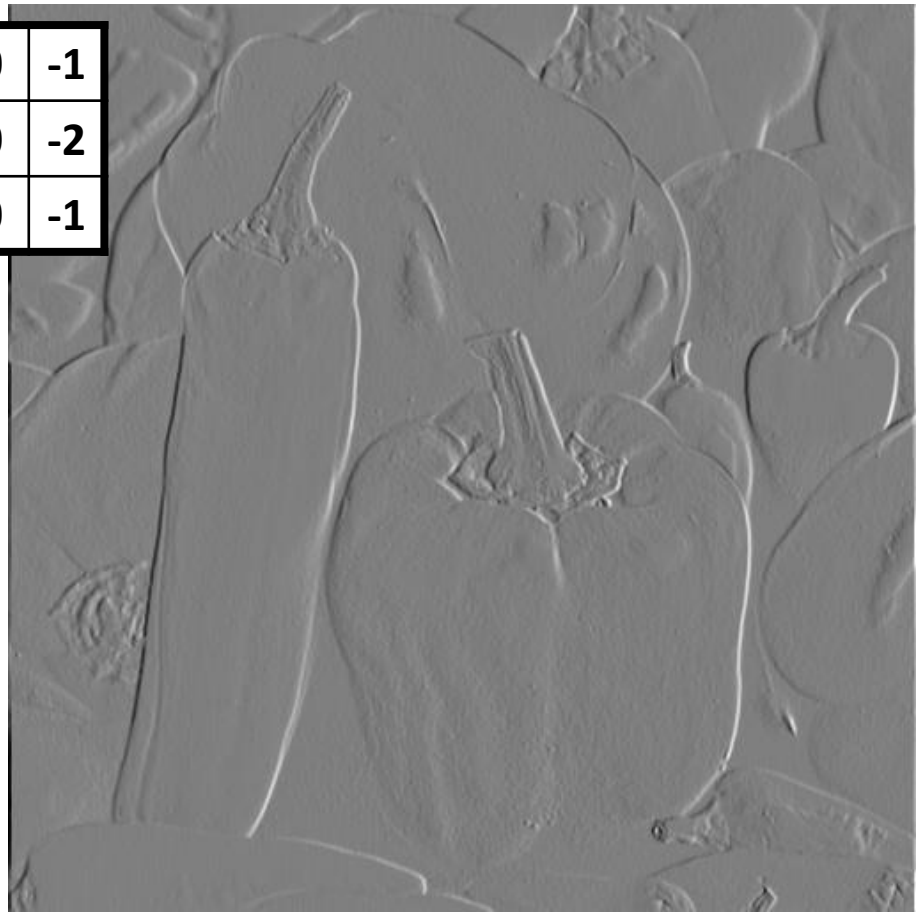

Common filters: Sobel

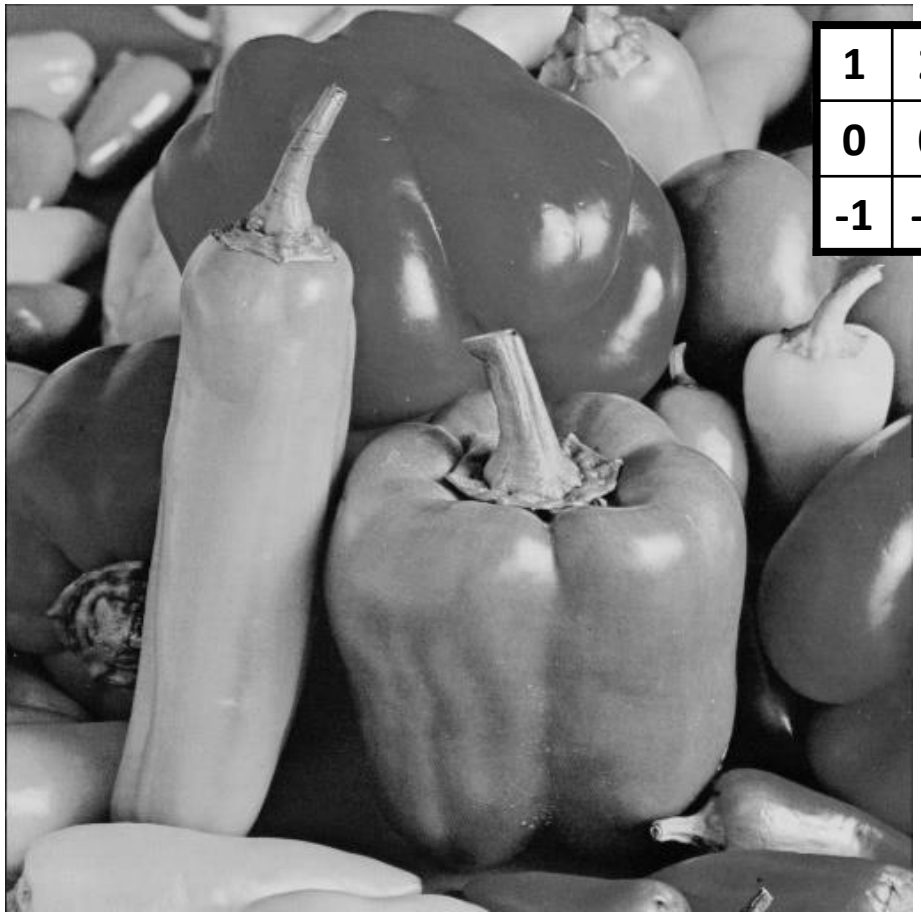

1	2	1
0	0	0
-1	-2	-1

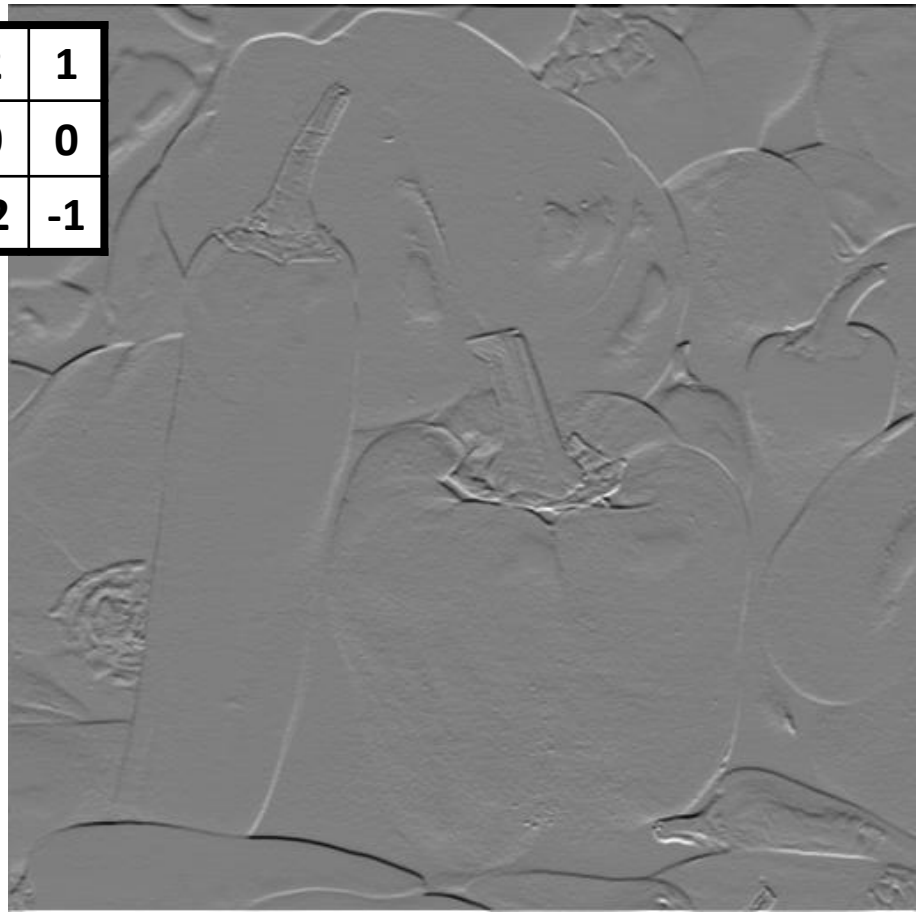

Designing filters

- How could you detect diagonal edges?

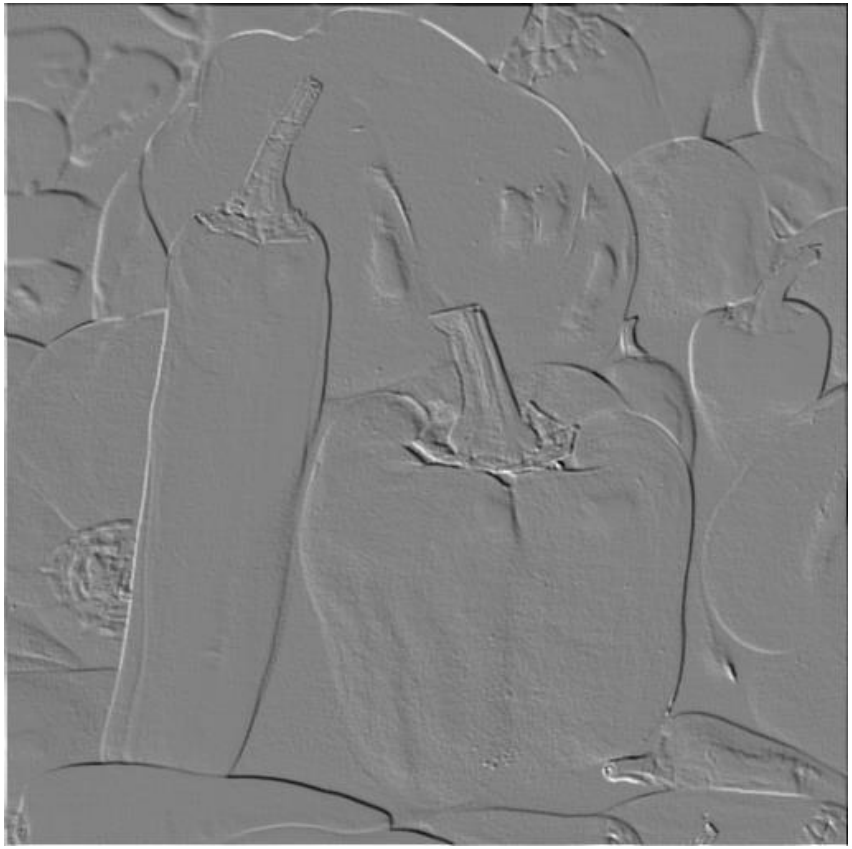

Designing filters

- How could you simulate (linear) motion blur?

Common filters

- Average/blur filters: average pixel values, blur the image
- Sharpening filters: subtract pixel from surround, increase fine detail
- Edge filters: compute difference between pixels, detect oriented edges in image

Filters in practice

Properties of linear filters

- Commutative: $f * h = h * f$
  - Theoretically, no difference between kernel and image
  - But most implementations do care about order
- Associative: $(f * h1) * h2 = f * (h1 * h2)$
  - Usually one option is faster than the other – allows for more efficient implementations
- Distributive over addition
  - $f * (h1 + h2) = (f * h1) + (f * h2)$
- Multiplication cancels out
  - $kf * h = f * kh = k(f * h)$

Efficient filtering

- Multiple filters: generally more efficient to combine 2D filters ($h1 * h2 * h3$…) and filter image just once

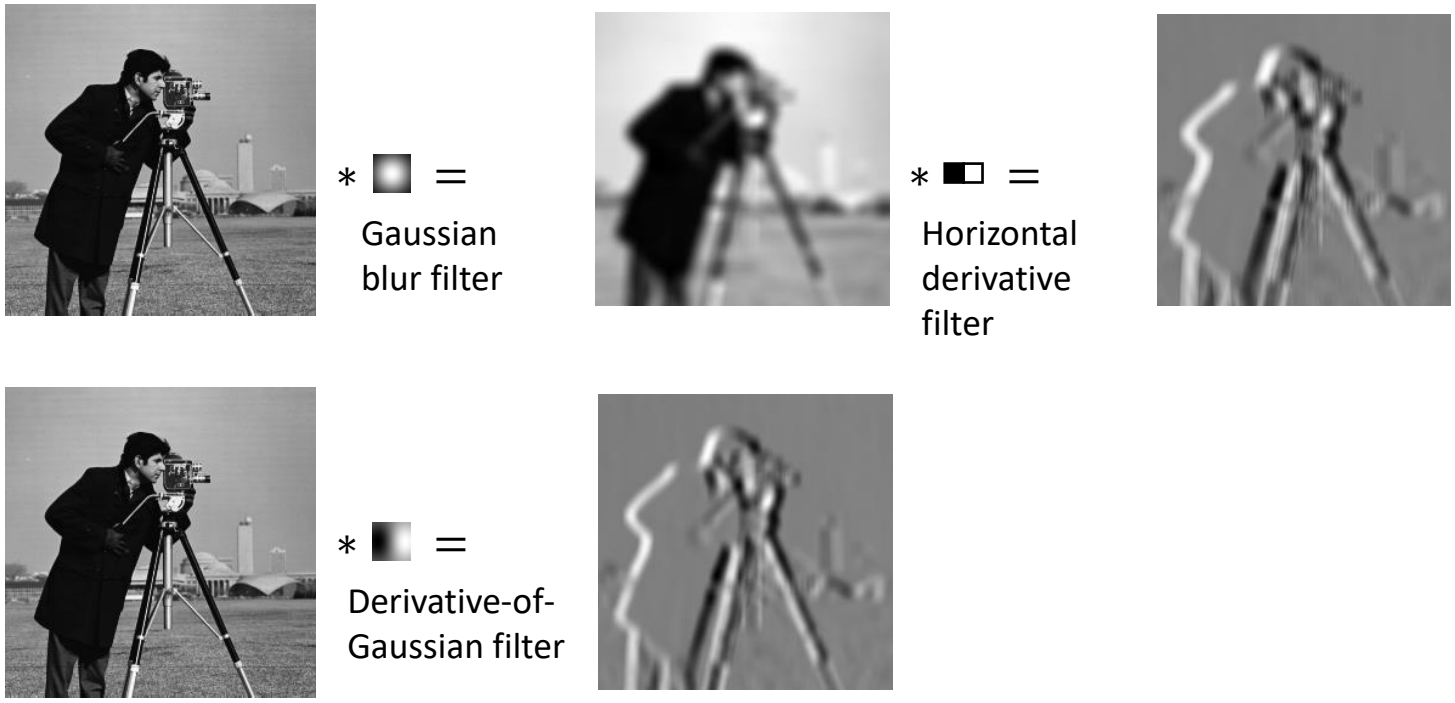

Efficient filtering

- Separable filters: generally more efficient to filter with two 1D filters than one 2D filter
- For example, the 2D Gaussian can be expressed as a product of two 1D Guassians (in x and y)

$$G_\sigma(x,y) = \frac{1}{2\pi\sigma^2}\exp^{-\frac{x^2+y^2}{2\sigma^2}}$$

$$= \left(\frac{1}{\sqrt{2\pi}\sigma}\exp^{-\frac{x^2}{2\sigma^2}}\right)\left(\frac{1}{\sqrt{2\pi}\sigma}\exp^{-\frac{y^2}{2\sigma^2}}\right)$$

Figure: D. Lowe

Separable filters

2D convolution
(center location only)

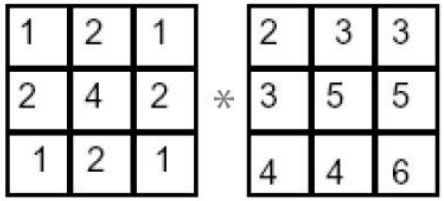

The filter factors into a product of 1D filters:

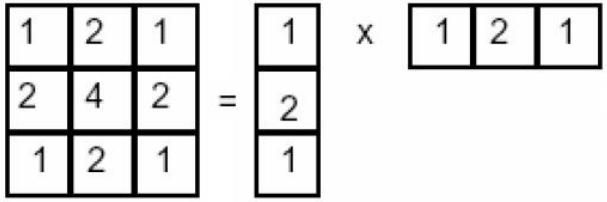

Perform convolution along rows:

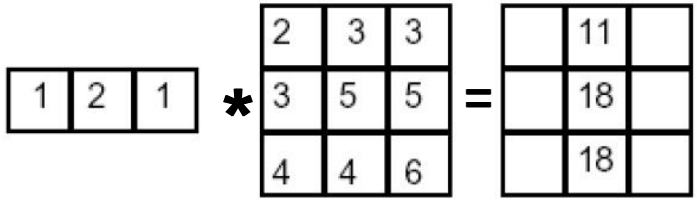

Followed by convolution along the remaining column:

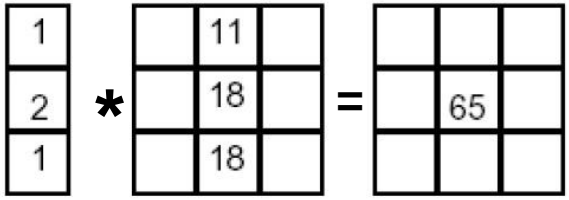

Slide: K. Grauman

Convolution output size

- Valid convolution: the output image is smaller than the input image
- Why?

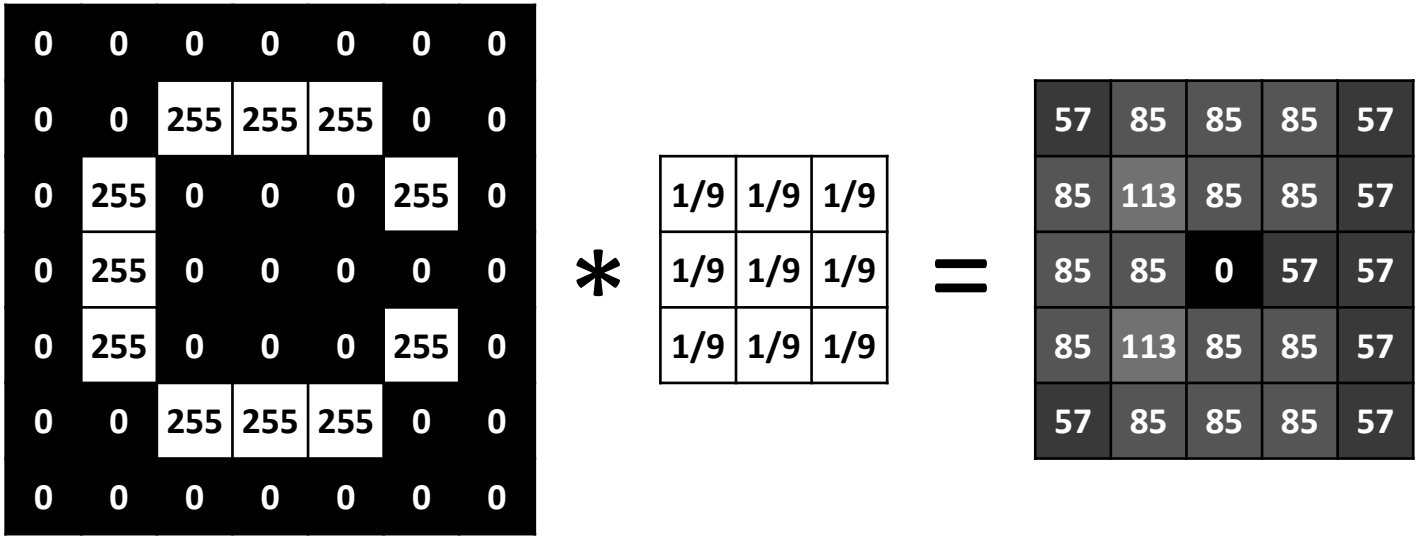

Border handling

- Pad with constant value

Filter: Gaussian blur

Border handling

- Wrap image

Filter: Gaussian blur

Border handling

- Clamp / replicate the border value

Filter: Gaussian blur

Border handling

- Reflect image

Filter: Gaussian blur

Practical considerations

- Think about how to implement filters efficiently
  - Images are big, so efficient filtering can save a lot of time!
- Think about how to handle borders
  - No one-size-fits-all solution
  - Wrap is ideal for tiling textures (but not photos)
  - Clamp/replicate tends to work well for photos

Summary

- Linear filters: first step of almost all computer vision systems
- Linear filters are just a first step – you can't build complex feature detectors from just linear filters